



Figure 1. Per-cohort and pooled AUC with 95% bootstrap CI under LODO cross-validation




Pooled LODO ROC Curves (n = 646)




Top 15 Species by SHAP Feature Importance (CRC vs. Healthy)


Random Forest




XGBoost



8/15 features shared in top 15 between RF and XGBoost




Alpha diversity (Shannon) by study_condition, per cohort




Species RF (Youden-J fixed): base-rate-adjusted PPV and NPV




PPV vs prevalence (sens=0.739, spec=0.698)






NPV vs prevalence (sens=0.739, spec=0.698)




Calibration mechanism — joint XGB pushes probabilities to 0/1, inflating reliability term



Reliability diagrams (10-bin equal-width) — pooled LODO predictions




Cohort composition by study condition




Confusion matrices — pooled LODO predictions


Species RF — Youden-J (t=0.485)
Sens=0.739 Spec=0.698 MCC=0.437




Species RF — Threshold = 0.5 (t=0.500)
Sens=0.697 Spec=0.728 MCC=0.425




Joint RF — Youden-J (t=0.519)
Sens=0.680 Spec=0.732 MCC=0.412




Joint RF — Threshold = 0.5 (t=0.500)
Sens=0.709 Spec=0.689 MCC=0.398




Joint XGB — Youden-J (t=0.533)
Sens=0.654 Spec=0.770 MCC=0.427



Joint XGB — Threshold = 0.5 (t=0.500)
Sens=0.668 Spec=0.750 MCC=0.419




Per-sample sequencing depth by cohort




F. nucleatum SHAP vs cohort sequencing depth


Spearman rho = -0.194, p = 0.591




PCoA (Bray-Curtis) coloured by cohort



PCoA (Bray-Curtis) coloured by study_condition




Species RF — per-cohort operating characteristics at pooled Youden-J ($t=0.485$)


SHAP vs permutation importance — species RF, CRC vs control

Top-20 overlap: 16/20


Permutation-importance rank (1 = largest AUC drop)




ROC — pooled LODO predictions




Precision-Recall — pooled LODO predictions



Sensitivity at clinically-meaningful specificity floors
(pooled LODO predictions, n = 1339)



Species RF — subgroup AUC (pooled LODO, 1000-iter bootstrap)



[illegible]